

B555: Homework 3

1. From Wasserman, Ch. 3 problems 1 (game with c dollars) , 4 (random walk) , 9 (Computer experiment) , 13 (Unif(0,1) and Unif(3,4)).
2. Let S be the collection of all of the binary strings no more than 4 digits long including the empty string e . Associate each string, $s \in S$ with a random variable, X_s such that

$$\begin{aligned} X_{s \circ 1} &= X_s + Y_{s \circ 1} \\ X_{s \circ 0} &= X_s + Y_{s \circ 0} \end{aligned}$$

where \circ denotes concatenation and X_e and $\{Y_s : s \in S\}$ are all mutually independent $N(0, 1)$ random variables.

- (a) Draw a directed acyclic graph labeling vertices with X_s random variables such that the graph represents the conditional independence structure of the probability distribution.
 - (b) Give a factorization of the joint density on the variables $\{X_s : s \in S\}$.
 - (c) What is the conditional expectation of X_e given all of the variables $\{X_s : s \in S, s \neq e\}$.
3. Suppose you observe univariate random variables $x_1, \dots, x_n \stackrel{\text{iid}}{\sim} N(\mu, \sigma^2)$ where the mean, μ , is known and the variance σ^2 is unknown. Find the maximum likelihood estimator for σ^2 , $\hat{\sigma}_{\text{MLE}}^2$.
 4. Suppose you observe d-dimensional vectors $x_1, \dots, x_n \stackrel{\text{iid}}{\sim} N(\mu, \Sigma)$ where the mean vector μ is unknown and the covariance matrix Σ is known. Find the maximum likelihood estimator for μ , $\hat{\mu}_{\text{MLE}}$.
 5. Fisher's famous iris data set measures sepal width, petal width, sepal length, and petal length for three different types of iris: Virginica, Versicolor, and Setosa. This problem requires you to build a Gaussian classifier using these data. To read the data set into R use

```
data(iris)
```

If, for instance you want to create the empirical mean and covariance matrix

$$\begin{aligned} \bar{x} &= \frac{1}{n} \sum_i x_i \\ S &= \frac{1}{n} \sum_i (x_i - \bar{x})(x_i - \bar{x})^t \end{aligned}$$

for the Versicolor class you could use

```
X = iris[iris[,5] == "versicolor",1:4]
mv = as.matrix(colMeans(X));
Sv = as.matrix(cov(X))
```

where the `as.matrix` cast is important. (The 5th column has the class labels and the first four columns are the feature values.) You can get the inverse and determinant of S with

```
Cv = solve(Sv)
dv = det(Sv)
```

- (a) Create a scatter plot using the first two features with different plot symbols for each iris type.
- (b) Using all of the data compute the empirical mean vectors and covariance matrices for the 3 classes. Use these to classify the data vectors using an optimal Bayes' classifier assuming the class conditionals are multivariate Gaussian. Show a scatter plot of the first two feature values as before, now coloring the correctly classified vectors in green and the incorrect ones in red.

6. In the “earthquakes.r” example (on the web page) we showed how our information about the Poisson rate, λ was updated sample-by-sample using the conjugate prior pair of

$$\lambda \sim \text{Gamma}(\theta, k)$$

$$x_1, \dots, x_n \stackrel{\text{iid}}{\sim} \text{Poisson}(\lambda)$$

Now assume that

$$\mu \sim \text{Normal}(\nu, \rho^2)$$

$$x_1, \dots, x_n \stackrel{\text{iid}}{\sim} \text{Normal}(\mu, \sigma^2)$$

with $\nu = 0, \rho^2 = 1,000,000, \sigma^2 = 1$.

- (a) In R, generate samples x_1, x_2, \dots , from the $\text{Normal}(10, 1)$ distribution and plot the conditional pdfs $p(\mu|x_1, \dots, x_n)$ for $n = 1 \dots 100$.
 - (b) Describe how the conditional densities change as you receive more and more data.
 - (c) What appears to be the limit of $E[\mu|x_1, \dots, x_n]$ as n increases?
7. In class we showed that for any x_1, \dots, x_n with sample mean \bar{x} , $\arg \min_c \sum_i (x_i - c)^2 = \bar{x}$. In a manner analogous to this argument show that for a continuous random variable with density $p(x)$ and expectation μ ,

$$\arg \min_c \int_{-\infty}^{\infty} (x - c)^2 p(x) dx = \mu$$

8. $X \sim \text{Beta}(\alpha, \beta)$ if X has pdf

$$p(x) = \begin{cases} \frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)} & 0 < x < 1 \\ 0 & \text{otherwise} \end{cases}$$

where $B(\alpha, \beta)$ is the constant necessary to make the pdf integrate to 1 (we don't need to know how to compute B in this problem). It may help to know that the mean of a $\text{Beta}(\alpha, \beta)$ distribution is $\frac{\alpha}{\alpha+\beta}$. Suppose $q \sim \text{Beta}(\alpha, \beta)$ and $X \sim \text{Binomial}(n, q)$.

- (a) What is the conditional density $p(q|x)$.
- (b) If $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \text{Binomial}(n, q)$ what is $p(q|x_1, \dots, x_n)$?
- (c) Suppose we want to estimate q with some number \hat{q} . if our estimate misses the true value of q by distance d we will pay $\$d^2$. How should \hat{q} be chosen so that we minimize our expected cost?